

SGN-41007 Pattern Recognition and Machine Learning

Exercise Set 2: January 18 - 20, 2017

Exercises consist of both pen&paper and computer assignments. Pen&paper questions are solved at home before exercises, while computer assignments are solved during exercise hours. The computer assignments are marked by text `python` and Pen&paper questions by text `pen&paper`

1. `pen&paper` Find a maximum likelihood estimator for one sample.

Suppose we measure one sample x for which we assume the model

$$p(x, \lambda) = \begin{cases} \lambda \exp(-\lambda x) & x \geq 0 \\ 0, & x < 0 \end{cases}$$

- Write and simplify the expression for $\ln p(x; \lambda)$. For simplicity, you can omit the case $x < 0$.
- Compute the derivative $\frac{\partial \ln p(x; \lambda)}{\partial \lambda}$.
- Solve $\frac{\partial \ln p(x; \lambda)}{\partial \lambda} = 0$ with respect to λ . This is the formula for maximum likelihood estimate of λ .

2. `pen&paper` Find a maximum likelihood estimator for many samples.

Suppose we measure samples $\mathbf{x} = (x[0], x[1], \dots, x[N-1])$ for which we assume the model

$$p(x[n], \lambda) = \begin{cases} \lambda \exp(-\lambda x[n]) & x[n] \geq 0 \\ 0, & x[n] < 0 \end{cases}$$

- Write and simplify the expression for $\ln p(\mathbf{x}; \lambda)$. You can assume the samples are independent and $p(\mathbf{x}; \lambda)$ is simply the product of individual probabilities: $p(\mathbf{x}; \lambda) = \prod p(x[n], \lambda)$.
- Compute the derivative $\frac{\partial \ln p(\mathbf{x}; \lambda)}{\partial \lambda}$.
- Solve $\frac{\partial \ln p(\mathbf{x}; \lambda)}{\partial \lambda} = 0$ with respect to λ . This is the formula for maximum likelihood estimate of λ .

3. `python` Same as last week Question 1, but without numpy.

Download the following file and extract the contents:

<http://www.cs.tut.fi/courses/SGN-41006/exercises/locationData.zip>

- Read the file into memory one line at a time (in a for loop). See similar example at the end of lecture slide set 1.

- b) Load the same data into another variable using `numpy.loadtxt`. Check that the contents of the two arrays are equal using `numpy.all` or `numpy.any`.

4. **python** *Implement two utility functions we need later.*

- a) Implement the following function:

```
p = gaussian(x, mu, sigma)
```

which returns the Gaussian density with parameters `mu` and `sigma` at point `x`. In mathematical notation, the function is:

$$p(x; \mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[-\frac{1}{2\sigma^2}(x - \mu)^2 \right]$$

- b) Implement the function:

```
p = log_gaussian(x, mu, sigma)
```

that returns $\ln p(x; \mu, \sigma)$. Do not use the previous function, because the straightforward solution `p = numpy.log(gaussian(x, mu, sigma))` would be numerically inaccurate. Instead, first manually simplify the expression

$$\ln p(x; \mu, \sigma) = \ln \left(\frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[-\frac{1}{2\sigma^2}(x - \mu)^2 \right] \right)$$

and write the corresponding code.

- c) Plot both functions for `mu = 0` and `sigma = 1` in the interval $x \in [-5, 5]$ to check that the code is correct. *Hint:* create a vector of `x`-coordinates using `numpy.linspace` and pass that through your functions and plot.

5. **python** *Estimate sinusoidal parameters.*

- a) Generate a 100-sample long synthetic test signal from the model:

$$x[n] = \sin(2\pi f_0 n) + w[n], \quad n = 0, 1, \dots, 99$$

with $f_0 = 0.017$ and $w[n] \sim \mathcal{N}(0, 0.25)$. Note that $w[n]$ is generated by `w = numpy.sqrt(0.25) * numpy.random.randn(100)`. Plot the result.

- b) Implement code from estimating the frequency of `x` using the maximum likelihood estimator:

$$\hat{f}_0 = \text{value of } f \text{ that maximizes } \left| \sum_{n=0}^{N-1} x(n) e^{-2\pi i f n} \right|.$$

Implementation is straightforward by noting that the sum expression is in fact a dot product:

$$\hat{f}_0 = \text{value of } f \text{ that maximizes } |\mathbf{x} \cdot \mathbf{e}|,$$

with $\mathbf{x} = (x_0, x_1, \dots, x_{N-1})$ and $\mathbf{e} = (e^{-2\pi i f \cdot 0}, e^{-2\pi i f \cdot 1}, \dots, e^{-2\pi i f \cdot (N-1)})$.

Use the following template and fill in the blanks.

```
scores = []
frequencies = []

for f in numpy.linspace(0, 0.5, 1000):

    # Create vector e. Assume data is in x.

    n = numpy.arange(100)
    z = # <compute -2*pi*i*f*n. Imaginary unit is 1j>
    e = numpy.exp(z)

    score = # <compute abs of dot product of x and e>
    scores.append(score)
    frequencies.append(f)

fHat = frequencies[np.argmax(scores)]
```

c) Run parts (a) and (b) a few times. Are the results close to true $f_0 = 0.017$?